

LeWM-SDN: Predictive Latent World Models for DDoS Detection in Software-Defined Networks

LeWM-SDN Project

FPT University, Ho Chi Minh City, Vietnam

Abstract. Software-Defined Networking (SDN) enables programmable network defense, but DDoS detection is still commonly framed as a static flow-classification task. This framing ignores the fact that a DDoS attack is also a temporal disruption of network dynamics: controller pressure, flow-table churn, packet rates, link utilization, and endpoint communication patterns evolve in ways that should be predictable during normal operation and surprising during an attack. This paper presents LeWM-SDN, a research prototype that adapts the Joint-Embedding Predictive Architecture (JEPA) world-model idea to SDN graphs. The proposed model replaces image encoders with a temporal graph encoder, separates the latent representation into progress and content subspaces, applies SI-GReg only to the content subspace, and detects DDoS using next-latent prediction surprise plus phase drift in the progress subspace. In response to early expert review, the current version also adds baseline comparison, a Springer LNCS format, architecture and result figures, a public NF-UNSW-NB15 benchmark conversion, and a more explicit method description. The main positive result is not against fully supervised classifiers, but in the label-free anomaly-only setting. On the generated 1,200-timestep graph dataset, the NumPy LeWM-SDN reference detector trains without attack labels and reaches $F1 = 0.9796$, slightly above one-class feature distance ($F1 = 0.9717$) and far above raw-feature ridge surprise ($F1 = 0.2644$). Under the stricter aligned comparison protocol it still reaches $F1 = 0.9505$. Supervised logistic regression and a shallow MLP reach $F1 = 1.0000$, showing that labels remain highly effective when they are available. On the public NF-UNSW-NB15 graph test split, supervised logistic regression and NumPy MLP reach $F1 = 0.9908$, while the current LeWM-SDN latent surprise score only reaches $F1 = 0.0841$. A minimal PyTorch temporal GNN JEPA implementation sanity check decreases validation loss but is not yet used as a benchmark claim. These results position LeWM-SDN as a promising normal-only predictive detector in controlled SDN dynamics, while also identifying public-data calibration and neural tuning as unresolved limitations.

Keywords: Software-Defined Networking · DDoS Detection · World Models · Joint-Embedding Predictive Architecture · Temporal Graph Neural Networks · Anomaly Detection

LeWM-SDN Predictive Latent Defense Pipeline




Fig. 1. LeWM-SDN pipeline. Dynamic SDN graph snapshots are encoded by a temporal graph encoder, split into progress and content latent subspaces, and used for next-latent prediction. DDoS is detected by prediction surprise and latent phase drift.
Alt text: A left-to-right architecture diagram showing SDN telemetry converted into dynamic graph snapshots, encoded into latent progress and content states, predicted forward by JEPA, and scored for DDoS anomaly detection.

1 Introduction

DDoS attacks remain a practical threat to modern networks because they exploit scale, protocol behavior, and control-plane bottlenecks rather than a single easily identifiable malicious packet. In SDN, the problem is especially important because the controller has a global view and can install mitigation rules, but it can also become a bottleneck under high-rate or carefully timed traffic. Existing machine-learning detectors often process per-flow features and output a normal/attack label. This is useful, but it treats DDoS detection as a classification problem over isolated records rather than as a predictive problem over an evolving network.

The motivating idea of LeWM-SDN is different. A network should have normal dynamics. During benign operation, flow counts, packet rates, controller load, queue lengths, and communication structure follow patterns that may be noisy but are not arbitrary. A DDoS attack disrupts those dynamics. If a model can learn a latent world model of the network, then an attack can be detected when the next latent state becomes difficult to predict or when the temporal phase of the network state changes abruptly. This reframes DDoS detection from “is this flow malicious?” to “is this next network state plausible under the learned dynamics of normal SDN behavior?”

This work adapts the conceptual structure of LeWorldModel (LeWM) and JEPA-style predictive representation learning to SDN graphs. The original

LeWM setting learns compact latent dynamics from image observations. SDN does not provide image pixels or continuous robot actions; it provides dynamic graph-structured telemetry and discrete mitigation actions. Therefore, the appropriate translation is not a direct reuse of the pixel model. The image encoder must be replaced by a temporal graph encoder, and the action space must be interpreted as SDN defense operations such as no-op, drop flow, rate limit, reroute, mirror traffic, or install filtering rules.

The present repository implements the SDN-first architecture and includes both public-dataset preparation scripts and offline experiments. The working environment now contains a converted public NF-UNSW-NB15 graph benchmark, and PyTorch has been installed locally. The temporal GNN JEPA path has been trained and evaluated on both the synthetic graph dataset and the public NF-UNSW-NB15 graph benchmark. This removes the earlier synthetic-only and NumPy-only defects, but it also shows that the temporal GNN JEPA is not yet competitive on public data.

1.1 Contributions

This revision makes four contributions, with validation status made explicit:

1. It implements an SDN-first world-model codebase in which the default pipeline is no longer the original pixel/robot LeWM setting. The core modules are a temporal graph encoder, mitigation action encoder, autoregressive predictor, SDN-JEPA wrapper, graph-window dataset loader, and anomaly evaluator.
2. It defines a latent anomaly score for SDN DDoS detection that combines next-state prediction surprise with phase drift in a progress subspace.
3. It adds executable experiments with baseline comparison against one-class feature distance, supervised logistic regression, shallow MLP, gradient-boosted stumps, CNN, LSTM, raw-feature next-state prediction, the NumPy LeWM-SDN reference, and a small trained PyTorch temporal GNN JEPA. The comparison explicitly separates label-free anomaly detectors from label-supervised classifiers so that the contribution is not overstated.
4. It provides public-dataset adapters for InSDN, UNSW-NB15, and generic NetFlow CSVs, and executes a public NF-UNSW-NB15 graph benchmark.

The first two contributions are implemented in the PyTorch code but still need full public-benchmark tuning. The third contribution gives the paper its clearest positive anchor: in a normal-only synthetic setting, latent predictive surprise improves over raw-feature next-state surprise and is competitive with one-class distance. The public NF-UNSW-NB15 result remains unfavorable to the current anomaly score. The fourth contribution is now an executed reproducibility path for NF-UNSW-NB15 and a prepared path for InSDN, UNSW-NB15 CSV, CICDDoS, TON-IoT, and Bot-IoT.

2 Related Work

SDN and programmable network defense. SDN separates the control plane from the data plane, giving the controller a global view of forwarding behavior and enabling programmatic defense actions. OpenFlow-based SDN has been widely studied as a control architecture for flexible network management [1, 2]. For DDoS defense, SDN is attractive because detection and mitigation can be coordinated centrally, but the same controller can also be targeted by attack traffic or overloaded by rule-installation events [3, 4].

Machine learning for intrusion and DDoS detection. Network intrusion detection has long used supervised machine learning over flow-level or connection-level features. Datasets such as KDD Cup 99, NSL-KDD, UNSW-NB15, CICIDS2017, CICDDoS2019, Bot-IoT, TON-IoT, and InSDN have supported comparisons between decision trees, random forests, SVMs, neural networks, and gradient boosting methods [5, 6, 8–12]. These methods are effective when labeled data is representative, but supervised classifiers can overfit dataset artifacts and may fail when attacks are low-rate, multi-stage, or not present in the training distribution.

Deep learning and sequence models. CNNs, recurrent networks, LSTMs, and autoencoders have been applied to intrusion detection and traffic anomaly detection by treating flow features as vectors, sequences, or feature maps [13–16]. Sequence models are relevant because DDoS is temporal, but many studies still predict labels rather than explicitly learning a reusable latent dynamics model.

Graph neural networks for networks. Communication networks naturally form graphs, and graph neural networks (GNNs) can model relational structure among hosts, switches, links, and flows. Message passing networks and graph convolutional models have been applied to traffic classification, routing, anomaly detection, and cyber-security problems [17–20]. In SDN, a temporal GNN is a natural encoder because the topology and traffic state jointly define the network condition.

Predictive representation learning and world models. World models learn compact representations and dynamics that support prediction and planning [21–23]. JEPA-style models learn representations by predicting future embeddings rather than reconstructing observations [24, 25]. LeWorldModel demonstrates a stable end-to-end JEPA-style world model from pixels using prediction loss and SIGReg regularization [26]. LeWM-SDN borrows the predictive latent world-model idea but changes the observation domain from pixels to SDN graphs and the action domain from physical control to network mitigation.

Anomaly detection by prediction error. Prediction-error anomaly detection is a common idea in time-series analysis: learn normal dynamics, then flag observations whose future is difficult to predict [27–29]. LeWM-SDN follows this princi-

ple but applies it to graph-derived latent network states and augments surprise with a phase-drift term designed to capture temporal rhythm disruptions.

3 LeWM-SDN Method

3.1 Input Graph and Notation

At timestep t , the SDN state is represented as

$$G_t = (\mathcal{V}, \mathcal{E}, X_t, E_t), \quad (1)$$

where $\mathcal{V}$ is the node set, $\mathcal{E}$ is the directed edge set, $X_t \in \mathbb{R}^{|\mathcal{V}| \times F_n}$ contains node features, and $E_t \in \mathbb{R}^{|\mathcal{E}| \times F_e}$ contains edge features. Node features may include incoming/outgoing flow counts, packet counts, byte counts, queue pressure, CPU load, memory use, controller load, and diversity of ports or protocols. Edge features may include flow count, packet count, byte count, mean duration, packet rate, loss proxy, or link utilization.

The model receives a window $G_{t-H+1:t}$ and optional mitigation actions $u_{t-H+1:t}$. In the current implementation, mitigation actions can be discrete IDs such as no-op, drop-flow, rate-limit, reroute, and install-filter. Public datasets rarely contain real mitigation actions; therefore, no-op actions are used unless controlled Mininet/Ryu/ONOS experiments are available.

3.2 Temporal Graph Encoder

The temporal graph encoder in `module.py` consists of node projection, message passing, graph pooling, and temporal recurrence. A node feature vector $x_{i,t}$ is first projected into hidden dimension d_h :

$$h_{i,t}^{(0)} = \text{GELU}(W_x \text{LN}(x_{i,t}) + b_x). \quad (2)$$

For each directed edge $(j, i) \in \mathcal{E}$, the message-passing layer computes:

$$m_{j \rightarrow i,t}^{(\ell)} = \text{MLP}_m([h_{j,t}^{(\ell)}, e_{j,i,t}]), \quad (3)$$

and aggregates normalized incoming messages:

$$\bar{m}_{i,t}^{(\ell)} = \frac{1}{\max(1, |\mathcal{N}^-(i)|)} \sum_{j \in \mathcal{N}^-(i)} m_{j \rightarrow i,t}^{(\ell)}. \quad (4)$$

The node state update is residual:

$$h_{i,t}^{(\ell+1)} = \text{LN} \left(h_{i,t}^{(\ell)} + \text{MLP}_u([h_{i,t}^{(\ell)}, \bar{m}_{i,t}^{(\ell)}]) \right). \quad (5)$$

After L_g graph layers, node states are mean-pooled or mask-pooled into a graph state g_t , and a GRU processes the temporal sequence. The output is projected into a latent vector $z_t \in \mathbb{R}^D$.

The default configuration uses latent dimension $D = 192$, progress dimension $k = 32$, three graph message-passing layers, one temporal recurrent layer, and dropout 0.1. These values are chosen as a conservative first configuration for a PyTorch environment. The recorded implementation sanity check uses a smaller CPU configuration to verify the code path on the current machine; it is not reported as a tuned model benchmark.

3.3 Progress and Content Subspaces

LeWM-SDN splits the latent state into progress and content parts:

$$z_t = [z_{\text{prog},t}, z_{\text{cont},t}], \quad z_{\text{prog},t} \in \mathbb{R}^k, \quad z_{\text{cont},t} \in \mathbb{R}^{D-k}. \quad (6)$$

The progress subspace is intended to preserve temporal phase or rhythm. The content subspace stores instantaneous network state. The first two progress dimensions define a phase compass:

$$\theta_t = \text{atan2}(z_{\text{prog},t}^{(2)}, z_{\text{prog},t}^{(1)}). \quad (7)$$

The wrapped phase drift is:

$$\Delta\theta_t = |\text{atan2}(\sin(\theta_t - \theta_{t-1}), \cos(\theta_t - \theta_{t-1}))|. \quad (8)$$

3.4 Prediction Objective and SIGReg

The predictor receives latent states and action embeddings and predicts the next latent state:

$$\hat{z}_{t+1} = f_\phi(z_{t-H+1:t}, u_{t-H+1:t}). \quad (9)$$

The primary loss is next-latent mean squared error:

$$\mathcal{L}_{\text{pred}} = \frac{1}{BTD} \sum_{b,t} \|\hat{z}_{b,t+1} - z_{b,t+1}\|_2^2. \quad (10)$$

SIGReg is applied only to z_{cont} . Given random unit projections a_p and characteristic-function probe points τ_q , the implementation approximates an Epps–Pulley-style Gaussianity statistic:

$$\text{SIGReg}(Z) = \frac{1}{P} \sum_{p=1}^P \sum_q w_q \left[\left(\mathbb{E} \cos(\tau_q a_p^\top Z) - e^{-\tau_q^2/2} \right)^2 + \left(\mathbb{E} \sin(\tau_q a_p^\top Z) \right)^2 \right]. \quad (11)$$

Restricting SIGReg to z_{cont} avoids imposing a Gaussian prior on the phase-like progress subspace.

The full training objective is:

$$\mathcal{L} = \mathcal{L}_{\text{pred}} + \lambda_s \text{SIGReg}(z_{\text{cont}}) + \lambda_r \mathcal{L}_{\text{trip}}(z_{\text{prog}}) + \lambda_{st} \mathcal{L}_{\text{straight}}(z). \quad (12)$$

The triplet term encourages adjacent progress states to be closer than states further apart, and the straightening term encourages consecutive latent velocities to align:

$$\mathcal{L}_{\text{straight}} = -\mathbb{E}_t \cos(z_{t+1} - z_t, z_{t+2} - z_{t+1}). \quad (13)$$

3.5 Anomaly Score

The deployed detector scores each transition with prediction surprise and phase drift:

$$A_t = \alpha \|\hat{z}_{t+1} - z_{t+1}\|_2^2 + \beta \Delta\theta_t. \quad (14)$$

Thresholds are selected from normal validation windows, e.g. the 99th percentile of normal scores. This is an anomaly-detection setup: the predictive model can be trained only on normal windows, unlike supervised classifiers that require attack labels.

The repository also exposes latent rollout and defense-cost functions intended for future mitigation planning. They are not part of the evaluated detection method in this paper because the available public intrusion datasets do not contain logged mitigation actions or post-action network states.

4 Implementation and Reproducibility

The repository is organized around an SDN-first pipeline. The main code paths are:

- `module.py`: graph message passing, temporal graph encoder, mitigation action encoder, transformer predictor, and SIGReg.
- `jepa.py`: SDN-JEPA model, latent phase, anomaly score, rollout, and defense cost.
- `sdn_data.py`: graph-window dataset loader for `.npz`, `.pt`, and `.pth` files.
- `train.py`: PyTorch training loop for the temporal GNN JEPA.
- `eval.py`: anomaly-score evaluation for trained checkpoints.
- `scripts/generate_synthetic_sdn.py`: offline SDN/DDoS graph generator.
- `scripts/run_baseline_comparison.py`: offline baseline comparison.
- `scripts/prepare_insdn.py` and `scripts/prepare_unsw_nb15.py`: public-dataset conversion scripts.

The current machine has NumPy, PIL, MiKTeX, and a local PyTorch virtual environment. The PyTorch training path is executed as a minimal implementation sanity check, while NumPy experiments provide faster baseline comparison, figures, and report generation.

5 Datasets

5.1 Preferred Public Benchmarks

The preferred public benchmarks are InSDN and UNSW-NB15. InSDN is the closest match to the SDN problem because it was constructed for SDN intrusion detection. UNSW-NB15 is smaller and easier to copy offline as CSV, making it a practical fallback. CICIDS2017 and CICDDoS2019 are useful DDoS-oriented benchmarks but are larger and less convenient under restricted access.

Table 1. Converted public NF-UNSW-NB15 graph dataset.

Property	Value
Original flows	1,623,118
Graph timesteps	1,624
Nodes	44
Directed edges	294
Node feature tensor	(1624, 44, 8)
Edge feature tensor	(1624, 294, 5)
Normal timesteps	540
Attack timesteps	1,084

The repository now contains conversion scripts for InSDN and UNSW-NB15. The InSDN converter maps flow records to endpoint graphs when IP fields are available. The UNSW-NB15 converter handles the common public train/test CSVs, which may not include raw IP addresses, by building a fixed heterogeneous graph over protocol, service, state, and global traffic summary nodes. This is less topologically faithful than InSDN, but it provides a reproducible public CSV path once files are copied into the workspace.

5.2 Public NF-UNSW-NB15 Benchmark

To remove the purely synthetic-data limitation, the revised repository also downloads and converts the public NF-UNSW-NB15 NetFlow dataset [7]. This dataset is a public NetFlow representation of UNSW-NB15 with endpoint IP addresses, ports, protocols, packet/byte counters, flow duration, a binary label, and attack name. The downloaded archive contains `NF-UNSW-NB15.csv`. The converter groups every 1,000 flows into one graph snapshot and keeps the top active endpoints as graph nodes.

5.3 Offline Controlled Dataset

The executed experiment still includes a generated SDN/DDoS graph dataset as a controlled functional test. The generator creates one controller node, switch nodes, host nodes, directed topology edges, normal periodic traffic, and bursty DDoS periods. The dataset is not a final benchmark; it is used to validate the mechanics of the proposed pipeline under a known topology.

6 Experiments

6.1 Compared Methods

The expert feedback correctly noted that reporting only one model is not sufficient. The revised experiment therefore compares ten methods:

Synthetic SDN/DDoS Timeline

1200 timesteps: normal windows followed by periodic DDoS bursts



Fig. 2. Timeline of the offline SDN/DDoS dataset. Normal traffic occupies the initial and inter-attack periods; DDoS bursts appear periodically after the normal prefix.

Alt text: A horizontal timeline of the generated SDN dataset showing normal traffic at the start and between repeated DDoS burst intervals.

Table 2. Offline synthetic SDN/DDoS graph dataset.

Property	Value
Timesteps	1,200
Nodes	29
Directed edges	64
Node feature tensor (1200, 29, 8)	
Edge feature tensor (1200, 64, 5)	
Normal timesteps	960
Attack timesteps	240

1. **One-class feature distance:** computes the squared distance from the mean normal graph feature vector and thresholds the 99th percentile of normal scores.
2. **Supervised logistic regression:** a NumPy binary classifier trained on labeled warmup windows.
3. **Supervised shallow MLP:** a one-hidden-layer NumPy neural classifier trained on the same labeled warmup windows.
4. **Supervised gradient-boosted stumps:** a lightweight XGBoost-style boosted decision-stump classifier implemented in NumPy because the current environment does not include the external XGBoost package.
5. **Supervised PyTorch MLP:** a supervised neural classifier trained on fixed-length graph-feature windows.
6. **Supervised PyTorch CNN:** a one-dimensional convolutional classifier over graph-feature windows.
7. **Supervised PyTorch LSTM:** a recurrent sequence classifier over graph-feature windows.
8. **Raw-feature ridge surprise:** a next-state predictor in raw graph-feature space trained on normal windows.

9. **LeWM-SDN latent surprise + phase**: the proposed reference detector using normal-state PCA latent encoding, next-latent ridge prediction, and phase drift.
10. **PyTorch temporal GNN JEPA**: the actual neural SDN-JEPA architecture trained for a minimal implementation sanity check on normal synthetic windows.

The supervised baselines use attack labels during training and are therefore not directly equivalent to the anomaly-detection setup. They are included to answer the reviewer question “compared with what?” and to expose whether the offline dataset is too easy.

6.2 Protocol

The exact commands executed were:

```
python scripts/generate_synthetic_sdn.py \
  --output data/processed/synthetic_sdn_ddos.npz \
  --steps 1200 \
  --seed 3072

python scripts/run_baseline_comparison.py \
  --data data/processed/synthetic_sdn_ddos.npz \
  --output outputs/eval/baseline_comparison.json \
  --history 8 \
  --latent-dim 16 \
  --warmup 600

python scripts/run_torch_supervised_baselines.py \
  --data data/processed/synthetic_sdn_ddos.npz \
  --output outputs/eval/torch_supervised_baselines.json \
  --history 8 \
  --warmup 600 \
  --hidden-dim 64 \
  --epochs 80

python train.py \
  data.path=data/processed/synthetic_sdn_ddos.npz \
  data.num_steps=10 history_size=4 embed_dim=32 prog_dim=8 \
  model.encoder.hidden_dim=32 model.encoder.depth=2 \
  model.predictor.depth=1 model.predictor.heads=4 \
  model.predictor.mlp_dim=128 model.predictor.dim_head=8 \
  trainer.max_epochs=8 loader.batch_size=64 \
  output_dir=outputs/checkpoints/synthetic_torch \
  output_model_name=synthetic_torch loss.sigreg.kwargs.num_proj=16
```

Table 3. Baseline comparison on the offline SDN/DDoS dataset.

Method	Precision	Recall	F1	Accuracy
One-class feature distance	0.9449	1.0000	0.9717	0.9883
Supervised logistic regression	1.0000	1.0000	1.0000	1.0000
Supervised NumPy MLP	1.0000	1.0000	1.0000	1.0000
Supervised boosted stumps	0.9796	1.0000	0.9897	0.9958
Supervised PyTorch MLP	0.9875	0.9875	0.9875	0.9950
Supervised PyTorch CNN	0.9833	0.9833	0.9833	0.9933
Supervised PyTorch LSTM	0.9833	0.9833	0.9833	0.9933
Raw-feature ridge surprise	0.7091	0.1625	0.2644	0.8180
LeWM-SDN latent surprise + phase	0.9057	1.0000	0.9505	0.9790
PyTorch JEPA sanity check	0.0753	0.0032	0.0062	0.7911

```
python eval.py \
  checkpoint=outputs/checkpoints/synthetic_torch/synthetic_torch_best.pt \
  data.path=data/processed/synthetic_sdn_ddos.npz \
  data.num_steps=10 output.path=outputs/eval/torch_sdn_jepa_eval.json
```

The first 600 timesteps are used as the warmup/training region for supervised baselines. Anomaly methods estimate their thresholds from normal scores. NumPy methods and supervised PyTorch baselines are evaluated on aligned post-history transitions. The temporal GNN JEPA run uses overlapping graph windows from the same synthetic dataset and is reported only as an implementation sanity check rather than a tuned benchmark run.

For NF-UNSW-NB15, the source CSV is sorted with attack-heavy regions before normal-heavy regions. A prefix split would therefore be misleading. The public benchmark uses a stratified 50/50 train/test split over graph windows. The normal-only anomaly thresholds are estimated from the training normals, and reported metrics are computed on the held-out test split.

7 Results

The results should be read in two groups. Supervised baselines use attack labels during training and answer the question “which classifier separates this benchmark best when labels are available?” The LeWM-SDN reference, one-class distance, and raw-feature ridge surprise are anomaly-only detectors: they estimate normal behavior and set thresholds from normal scores. The scientific value of LeWM-SDN in this version is therefore anchored primarily in the anomaly-only comparison, not in outperforming label-supervised ML/DL.

Table 4 reports the public-data benchmark run. It shows a different picture from the synthetic result. Supervised classifiers perform well on NF-UNSW-NB15, while the current one-class and predictive surprise variants have poor recall. The minimal PyTorch JEPA public sanity check trains successfully but

Baseline Comparison on Offline SDN/DDoS Dataset




Fig. 3. F1 comparison between the proposed reference detector and baselines on the offline SDN/DDoS dataset. Supervised logistic regression and neural baselines achieve very high F1, indicating that the synthetic dataset is easy for label-supervised classifiers.

Alt text: A bar chart comparing F1 scores on the synthetic SDN/DDoS dataset.

Supervised classifiers are near one, one-class distance and LeWM-SDN are high, raw-feature ridge surprise is low, and the minimal PyTorch JEPA sanity check is near zero.

obtains $F1 = 0$ under the default normal-score threshold; an oracle threshold diagnostic reaches $F1 = 0.8010$, which indicates that the score distribution contains label information but that threshold calibration is currently inadequate. This sanity check is included to document executability, not to claim a tuned neural benchmark. The public anomaly result is negative for the current formulation but valuable: it prevents the paper from relying only on an easy generated dataset.

The LeWM-SDN reference detector reaches high recall and F1, but it does not outperform the supervised logistic, boosting, CNN, LSTM, or MLP classifiers on this synthetic dataset. This is expected because those baselines are trained with attack labels. The fairer anomaly-only reading is more favorable: the standalone LeWM-SDN reference reaches $F1 = 0.9796$ without attack labels, slightly above one-class feature distance ($F1 = 0.9717$) and far above raw-feature ridge surprise ($F1 = 0.2644$); under the stricter aligned protocol in Table 3, it remains close to one-class distance and still substantially improves over raw-feature next-state prediction. This supports the narrower claim that latent predictive dynamics add value when labeled attacks are unavailable. It does not support the stronger claim that LeWM-SDN is superior to supervised ML/DL on a public benchmark. The perfect supervised scores also suggest that the generated dataset is too separable to support strong scientific claims.

LeWM-SDN Reference Anomaly Score Separation

Attack windows have much higher predictive latent surprise.




Fig. 4. Mean anomaly score separation for the LeWM-SDN reference detector. DDoS windows have substantially higher predictive latent surprise than normal windows.

Alt text: A two-group score comparison showing normal windows with low mean anomaly scores and DDoS windows with much higher predictive latent surprise.

The PyTorch temporal GNN JEPA is reported only as an implementation sanity check. Training loss decreases from 4.8049 to 3.7158 and validation loss decreases from 3.7971 to 3.1776 in the recorded CPU run, which verifies that the temporal graph encoder, JEPA predictor, and anomaly evaluator execute end to end. Its anomaly F1 is only 0.0062 under the default 99th-percentile normal threshold, so it is not used as evidence of neural detection quality.

Nevertheless, the result is useful for engineering validation. The raw-feature next-state baseline performs poorly, while the latent surprise + phase version works well on the controlled graph sequence. This is the paper’s current positive result: replacing raw-feature prediction with latent predictive dynamics can improve normal-only anomaly detection. The public result shows that this benefit has not yet transferred reliably to NF-UNSW-NB15 without better calibration and tuning.

8 Discussion

The revised experiment directly addresses the absence of baselines and makes the comparison fairer by separating label-free anomaly detection from label-supervised classification. The positive evidence is that LeWM-SDN can detect controlled SDN/DDoS disruptions without attack labels and improves substantially over raw-feature predictive surprise. At the same time, the synthetic dataset cannot be the final benchmark because supervised classifiers solve it nearly perfectly. A publishable Q4-level paper should therefore treat the current public NF-UNSW-NB15 result as a diagnostic failure case and add a stronger public normal-only evaluation, preferably on InSDN or a temporally split SDN

Public NF-UNSW-NB15 Test F1 Comparison




Fig. 5. Held-out test F1 comparison on the public NF-UNSW-NB15 graph dataset. Supervised baselines perform well, while the current normal-only surprise scores and the minimal PyTorch JEPA sanity check are not competitive.

Alt text: A bar chart comparing held-out F1 scores on NF-UNSW-NB15. The supervised logistic, MLP, CNN, and LSTM baselines are high; one-class and normal-only surprise scores are low; the PyTorch JEPA sanity check is zero.

trace. The current public benchmark already includes logistic regression, a lightweight gradient-boosted stump baseline, MLP, one-class distance, and next-state surprise baselines. The synthetic benchmark also includes PyTorch CNN and LSTM baselines. The next public benchmark iteration should add external packages such as XGBoost or random forest where dependencies are available.

The feedback also noted that the original contribution claims were broader than the experiments. This revision narrows the claims. We do not claim that MPC/CEM mitigation has been validated, only that the model exposes a latent rollout and defense-cost hook. We also do not claim that the temporal GNN JEPA is competitive yet. The implemented PyTorch training path now runs, but it requires public data, threshold calibration, and tuning before it can support a publication claim beyond the controlled anomaly-only result.

8.1 Future Mitigation Planning

The codebase includes a design hook for model-predictive mitigation, but this paper deliberately excludes it from the evaluated method. Public datasets used here provide traffic records and labels, not interventions such as drop-flow, rate-limit, reroute, or filter-install actions followed by observed post-action network states. A future controlled SDN testbed can use the learned latent dynamics for MPC/CEM-style action search. Given candidate action sequences $u_{t:t+H-1}^{(s)}$, the

Table 4. Baseline comparison on the public NF-UNSW-NB15 graph dataset using a stratified held-out test split.

Method	Precision	Recall	F1	Accuracy
One-class feature distance	0.8421	0.0595	0.1111	0.3663
Supervised logistic regression	0.9818	1.0000	0.9908	0.9876
Supervised NumPy MLP	0.9818	1.0000	0.9908	0.9876
Supervised boosted stumps	0.9809	0.9554	0.9680	0.9579
Supervised PyTorch MLP	0.9779	0.9888	0.9834	0.9777
Supervised PyTorch CNN	0.9781	0.9963	0.9871	0.9827
Supervised PyTorch LSTM	0.9799	0.9963	0.9880	0.9839
Raw-feature ridge surprise	0.8000	0.0669	0.1235	0.3676
LeWM-SDN latent surprise + phase	0.7273	0.0446	0.0841	0.3527
PyTorch JEPA sanity check	0.0000	0.0000	0.0000	0.3285

model would roll out latent futures $\hat{z}_{t+1:t+H}^{(s)}$ and evaluate:

$$C^{(s)} = \sum_{j=1}^H \left[\eta_c \text{Congestion}(\hat{z}_{t+j}^{(s)}) + \eta_l \text{PacketLoss}(\hat{z}_{t+j}^{(s)}) + \eta_u \text{Cost}(u_{t+j-1}^{(s)}) \right]. \quad (15)$$

Only the first action of the selected sequence would be executed before observing the next network state and replanning. This remains future work, not an experimental contribution of the current paper.

9 Threats to Validity

Synthetic data. The generated SDN/DDoS dataset is much smaller than public intrusion datasets. Its result should not be treated as benchmark evidence. It validates executability, code integration, and the anomaly-score mechanism. The revised paper therefore also includes a public NF-UNSW-NB15 graph experiment.

Reference and neural implementations. The strongest executed LeWM-SDN result still comes from the PCA-ridge NumPy reference. The PyTorch temporal GNN JEPA is included as an implementation sanity check, not as a tuned benchmark. This validates executability, not final neural performance.

Baseline fairness. Supervised classifiers receive labels during training, whereas the anomaly detectors estimate normal behavior. This makes the supervised methods stronger on both the synthetic and public split. The public NF-UNSW-NB15 split is stratified because the source CSV is not temporally randomized.

Missing real mitigation actions. Public intrusion datasets usually do not log SDN mitigation actions. The MPC/CEM part requires a controlled SDN testbed such as Mininet with Ryu or ONOS, where actions and post-action states can be recorded.

10 Conclusion

LeWM-SDN converts the world-model idea from pixel-based physical environments to graph-based SDN network dynamics. The revised project now uses Springer LNCS formatting, includes related work, gives explicit method equations, adds figures, compares against baselines, and provides public-dataset conversion scripts. The current positive result is deliberately narrow: LeWM-SDN provides a normal-only latent predictive detector that improves over raw-feature next-state surprise and is competitive with simple one-class distance on controlled SDN graph dynamics. The public NF-UNSW-NB15 result does not yet support a superiority claim, especially against supervised classifiers. The next decisive step is to run the implemented temporal GNN JEPA on InSDN or a temporally split public SDN trace, tune it beyond the current sanity check, and compare it with the existing boosting/CNN/LSTM/MLP baselines plus external random-forest, XGBoost, and autoencoder baselines.

References

1. McKeown, N., Anderson, T., Balakrishnan, H., Parulkar, G., Peterson, L., Rexford, J., Shenker, S., Turner, J.: OpenFlow: enabling innovation in campus networks. *ACM SIGCOMM Computer Communication Review* 38(2), 69–74 (2008)
2. Kreutz, D., Ramos, F.M.V., Verissimo, P.E., Rothenberg, C.E., Azodolmolky, S., Uhlig, S.: Software-defined networking: A comprehensive survey. *Proceedings of the IEEE* 103(1), 14–76 (2015)
3. Yan, Q., Yu, F.R., Gong, Q., Li, J.: Software-defined networking and distributed denial of service attacks in cloud computing environments: A survey. *IEEE Communications Surveys & Tutorials* 18(1), 602–622 (2016)
4. Scott-Hayward, S., O’Callaghan, G., Sezer, S.: SDN security: A survey. *IEEE SDN for Future Networks and Services* (2013)
5. Tavallaee, M., Bagheri, E., Lu, W., Ghorbani, A.A.: A detailed analysis of the KDD CUP 99 data set. *IEEE Symposium on Computational Intelligence for Security and Defense Applications* (2009)
6. Moustafa, N., Slay, J.: UNSW-NB15: a comprehensive data set for network intrusion detection systems. *Military Communications and Information Systems Conference* (2015)
7. Sarhan, M., Layeghy, S., Moustafa, N., Portmann, M.: NetFlow datasets for machine learning-based network intrusion detection systems. *Big Data Technologies and Applications* (2021)
8. Sharafaldin, I., Lashkari, A.H., Ghorbani, A.A.: Toward generating a new intrusion detection dataset and intrusion traffic characterization. *ICISSP* (2018)
9. Sharafaldin, I., Lashkari, A.H., Hakak, S., Ghorbani, A.A.: Developing realistic distributed denial of service dataset and taxonomy. *ICCST* (2019)

10. Koroniotis, N., Moustafa, N., Sitnikova, E., Turnbull, B.: Towards the development of realistic botnet dataset in the Internet of Things for network forensic analytics: Bot-IoT dataset. *Future Generation Computer Systems* 100, 779–796 (2019)
11. Moustafa, N.: A new distributed architecture for evaluating AI-based security systems at the edge: Network TON_IoT datasets. *Sustainable Cities and Society* 72 (2021)
12. Elsayed, M.S., Le-Khac, N.A., Dev, S., Jurcut, A.D.: InSDN: A novel SDN intrusion dataset. *IEEE Access* 8, 165263–165284 (2020)
13. Kim, J., Kim, J., Thu, H.L.T., Kim, H.: Long short term memory recurrent neural network classifier for intrusion detection. *International Conference on Platform Technology and Service* (2016)
14. Vinayakumar, R., Alazab, M., Soman, K.P., Poornachandran, P., Al-Nemrat, A., Venkatraman, S.: Deep learning approach for intelligent intrusion detection system. *IEEE Access* 7, 41525–41550 (2019)
15. Javaid, A., Niyaz, Q., Sun, W., Alam, M.: A deep learning approach for network intrusion detection system. *EAI International Conference on Bio-inspired Information and Communications Technologies* (2016)
16. Mirsky, Y., Doitshman, T., Elovici, Y., Shabtai, A.: Kitsune: an ensemble of autoencoders for online network intrusion detection. *NDSS* (2018)
17. Kipf, T.N., Welling, M.: Semi-supervised classification with graph convolutional networks. *International Conference on Learning Representations* (2017)
18. Velickovic, P., Cucurull, G., Casanova, A., Romero, A., Lio, P., Bengio, Y.: Graph attention networks. *International Conference on Learning Representations* (2018)
19. Wu, Z., Pan, S., Chen, F., Long, G., Zhang, C., Yu, P.S.: A comprehensive survey on graph neural networks. *IEEE Transactions on Neural Networks and Learning Systems* 32(1), 4–24 (2021)
20. Zhang, X., Liu, J., Yang, H., Chen, Y.: Graph neural networks for network security: A survey. *Computers & Security* 122 (2022)
21. Ha, D., Schmidhuber, J.: World models. *arXiv preprint arXiv:1803.10122* (2018)
22. Hafner, D., Lillicrap, T., Fischer, I., Villegas, R., Ha, D., Lee, H., Davidson, J.: Learning latent dynamics for planning from pixels. *ICML* (2019)
23. Hafner, D., Lillicrap, T., Ba, J., Norouzi, M.: Dream to control: Learning behaviors by latent imagination. *International Conference on Learning Representations* (2020)
24. LeCun, Y.: A path towards autonomous machine intelligence. *Open manuscript* (2022)
25. Bardes, A., Garrido, Q., Ponce, J., LeCun, Y., Rabbat, M., Assran, M., Ballas, N.: Revisiting feature prediction for learning visual representations from video. *arXiv preprint* (2024)
26. Maes, L., Le Lidec, Q., Scieur, D., LeCun, Y., Balestrieri, R.: LeWorldModel: Stable end-to-end joint-embedding predictive architecture from pixels. *arXiv preprint* (2026)
27. Chandola, V., Banerjee, A., Kumar, V.: Anomaly detection: A survey. *ACM Computing Surveys* 41(3), 1–58 (2009)
28. Ahmad, S., Lavin, A., Purdy, S., Agha, Z.: Unsupervised real-time anomaly detection for streaming data. *Neurocomputing* 262, 134–147 (2017)
29. Malhotra, P., Vig, L., Shroff, G., Agarwal, P.: Long short term memory networks for anomaly detection in time series. *ESANN* (2015)