

LeWM-SDN: Predictive Latent World Models for DDoS Detection in Software-Defined Networks

LeWM-SDN Project

FPT University, Ho Chi Minh City, Vietnam

Abstract. Software-Defined Networking (SDN) enables programmable network defense, but DDoS detection is still commonly framed as a static flow-classification task. This framing ignores the fact that a DDoS attack is also a temporal disruption of network dynamics: controller pressure, flow-table churn, packet rates, link utilization, and endpoint communication patterns evolve in ways that should be predictable during normal operation and surprising during an attack. This paper presents LeWM-SDN, a research prototype that adapts the Joint-Embedding Predictive Architecture (JEPA) world-model idea to SDN graphs. The proposed model replaces image encoders with a temporal graph encoder, separates the latent representation into progress and content subspaces, applies SIGReg only to the content subspace, and detects DDoS using next-latent prediction surprise plus phase drift in the progress subspace. In response to early expert review, the current version also adds baseline comparison, a Springer LNCS format, architecture and result figures, a public NF-UNSW-NB15 benchmark conversion, and a more explicit method description. On the generated 1,200-timestep graph dataset, the NumPy LeWM-SDN reference detector reaches $F1 = 0.9796$, while supervised logistic regression and a shallow MLP reach $F1 = 1.0000$. On the public NF-UNSW-NB15 graph test split, supervised logistic regression and NumPy MLP reach $F1 = 0.9908$, while the current LeWM-SDN latent surprise score only reaches $F1 = 0.0841$. Small 8-epoch PyTorch temporal GNN JEPA runs decrease validation loss on both synthetic and public graphs, but the public run obtains $F1 = 0$ under the default normal-score threshold. These results should be interpreted as honest prototype validation rather than a claim of superiority over ML/DL baselines.

Keywords: Software-Defined Networking · DDoS Detection · World Models · Joint-Embedding Predictive Architecture · Temporal Graph Neural Networks · Anomaly Detection

1 Introduction

DDoS attacks remain a practical threat to modern networks because they exploit scale, protocol behavior, and control-plane bottlenecks rather than a single easily identifiable malicious packet. In SDN, the problem is especially important because the controller has a global view and can install mitigation rules, but it can also become a bottleneck under high-rate or carefully timed traffic. Existing machine-learning detectors often process per-flow features and output a

LeWM-SDN Predictive Latent Defense Pipeline




Fig. 1. LeWM-SDN pipeline. Dynamic SDN graph snapshots are encoded by a temporal graph encoder, split into progress and content latent subspaces, and used for next-latent prediction. DDoS is detected by prediction surprise and latent phase drift.

normal/attack label. This is useful, but it treats DDoS detection as a classification problem over isolated records rather than as a predictive problem over an evolving network.

The motivating idea of LeWM-SDN is different. A network should have normal dynamics. During benign operation, flow counts, packet rates, controller load, queue lengths, and communication structure follow patterns that may be noisy but are not arbitrary. A DDoS attack disrupts those dynamics. If a model can learn a latent world model of the network, then an attack can be detected when the next latent state becomes difficult to predict or when the temporal phase of the network state changes abruptly. This reframes DDoS detection from “is this flow malicious?” to “is this next network state plausible under the learned dynamics of normal SDN behavior?”

This work adapts the conceptual structure of LeWorldModel (LeWM) and JEPA-style predictive representation learning to SDN graphs. The original LeWM setting learns compact latent dynamics from image observations. SDN does not provide image pixels or continuous robot actions; it provides dynamic graph-structured telemetry and discrete mitigation actions. Therefore, the appropriate translation is not a direct reuse of the pixel model. The image encoder must be replaced by a temporal graph encoder, and the action space must be interpreted as SDN defense operations such as no-op, drop flow, rate limit, reroute, mirror traffic, or install filtering rules.

The present repository implements the SDN-first architecture and includes both public-dataset preparation scripts and offline experiments. The working environment now contains a converted public NF-UNSW-NB15 graph benchmark, and PyTorch has been installed locally. The temporal GNN JEPA path has

been trained and evaluated on both the synthetic graph dataset and the public NF-UNSW-NB15 graph benchmark. This removes the earlier synthetic-only and NumPy-only defects, but it also shows that the temporal GNN JEPA is not yet competitive on public data.

1.1 Contributions

This revision makes four contributions, with validation status made explicit:

1. It implements an SDN-first world-model codebase in which the default pipeline is no longer the original pixel/robot LeWM setting. The core modules are a temporal graph encoder, mitigation action encoder, autoregressive predictor, SDN-JEPA wrapper, graph-window dataset loader, and anomaly evaluator.
2. It defines a latent anomaly score for SDN DDoS detection that combines next-state prediction surprise with phase drift in a progress subspace.
3. It adds executable experiments with baseline comparison against one-class feature distance, supervised logistic regression, shallow MLP, gradient-boosted stumps, CNN, LSTM, raw-feature next-state prediction, the NumPy LeWM-SDN reference, and a small trained PyTorch temporal GNN JEPA.
4. It provides public-dataset adapters for InSDN, UNSW-NB15, and generic NetFlow CSVs, and executes a public NF-UNSW-NB15 graph benchmark.

The first two contributions are implemented in the PyTorch code but still need full public-benchmark tuning. The third contribution is executed on both the synthetic graph and the public NF-UNSW-NB15 graph, although the public result is unfavorable to the current anomaly score. The fourth contribution is now an executed reproducibility path for NF-UNSW-NB15 and a prepared path for InSDN, UNSW-NB15 CSV, CICDDoS, TON-IoT, and Bot-IoT.

2 Related Work

SDN and programmable network defense. SDN separates the control plane from the data plane, giving the controller a global view of forwarding behavior and enabling programmatic defense actions. OpenFlow-based SDN has been widely studied as a control architecture for flexible network management [1, 2]. For DDoS defense, SDN is attractive because detection and mitigation can be coordinated centrally, but the same controller can also be targeted by attack traffic or overloaded by rule-installation events [3, 4].

Machine learning for intrusion and DDoS detection. Network intrusion detection has long used supervised machine learning over flow-level or connection-level features. Datasets such as KDD Cup 99, NSL-KDD, UNSW-NB15, CICIDS2017, CICDDoS2019, Bot-IoT, TON-IoT, and InSDN have supported comparisons

between decision trees, random forests, SVMs, neural networks, and gradient boosting methods [5, 6, 8–12]. These methods are effective when labeled data is representative, but supervised classifiers can overfit dataset artifacts and may fail when attacks are low-rate, multi-stage, or not present in the training distribution.

Deep learning and sequence models. CNNs, recurrent networks, LSTMs, and autoencoders have been applied to intrusion detection and traffic anomaly detection by treating flow features as vectors, sequences, or feature maps [13–16]. Sequence models are relevant because DDoS is temporal, but many studies still predict labels rather than explicitly learning a reusable latent dynamics model.

Graph neural networks for networks. Communication networks naturally form graphs, and graph neural networks (GNNs) can model relational structure among hosts, switches, links, and flows. Message passing networks and graph convolutional models have been applied to traffic classification, routing, anomaly detection, and cyber-security problems [17–20]. In SDN, a temporal GNN is a natural encoder because the topology and traffic state jointly define the network condition.

Predictive representation learning and world models. World models learn compact representations and dynamics that support prediction and planning [21–23]. JEPA-style models learn representations by predicting future embeddings rather than reconstructing observations [24, 25]. LeWorldModel demonstrates a stable end-to-end JEPA-style world model from pixels using prediction loss and SIGReg regularization [26]. LeWM-SDN borrows the predictive latent world-model idea but changes the observation domain from pixels to SDN graphs and the action domain from physical control to network mitigation.

Anomaly detection by prediction error. Prediction-error anomaly detection is a common idea in time-series analysis: learn normal dynamics, then flag observations whose future is difficult to predict [27–29]. LeWM-SDN follows this principle but applies it to graph-derived latent network states and augments surprise with a phase-drift term designed to capture temporal rhythm disruptions.

3 LeWM-SDN Method

3.1 Input Graph and Notation

At timestep t , the SDN state is represented as

$$G_t = (\mathcal{V}, \mathcal{E}, X_t, E_t), \quad (1)$$

where $\mathcal{V}$ is the node set, $\mathcal{E}$ is the directed edge set, $X_t \in \mathbb{R}^{|\mathcal{V}| \times F_n}$ contains node features, and $E_t \in \mathbb{R}^{|\mathcal{E}| \times F_e}$ contains edge features. Node features may include incoming/outgoing flow counts, packet counts, byte counts, queue pressure, CPU

load, memory use, controller load, and diversity of ports or protocols. Edge features may include flow count, packet count, byte count, mean duration, packet rate, loss proxy, or link utilization.

The model receives a window $G_{t-H+1:t}$ and optional mitigation actions $u_{t-H+1:t}$. In the current implementation, mitigation actions can be discrete IDs such as no-op, drop-flow, rate-limit, reroute, and install-filter. Public datasets rarely contain real mitigation actions; therefore, no-op actions are used unless controlled Mininet/Ryu/ONOS experiments are available.

3.2 Temporal Graph Encoder

The temporal graph encoder in `module.py` consists of node projection, message passing, graph pooling, and temporal recurrence. A node feature vector $x_{i,t}$ is first projected into hidden dimension d_h :

$$h_{i,t}^{(0)} = \text{GELU}(W_x \text{LN}(x_{i,t}) + b_x). \quad (2)$$

For each directed edge $(j, i) \in \mathcal{E}$, the message-passing layer computes:

$$m_{j \rightarrow i,t}^{(\ell)} = \text{MLP}_m([h_{j,t}^{(\ell)}, e_{j,i,t}]), \quad (3)$$

and aggregates normalized incoming messages:

$$\bar{m}_{i,t}^{(\ell)} = \frac{1}{\max(1, |\mathcal{N}^-(i)|)} \sum_{j \in \mathcal{N}^-(i)} m_{j \rightarrow i,t}^{(\ell)}. \quad (4)$$

The node state update is residual:

$$h_{i,t}^{(\ell+1)} = \text{LN} \left(h_{i,t}^{(\ell)} + \text{MLP}_u([h_{i,t}^{(\ell)}, \bar{m}_{i,t}^{(\ell)}]) \right). \quad (5)$$

After L_g graph layers, node states are mean-pooled or mask-pooled into a graph state g_t , and a GRU processes the temporal sequence. The output is projected into a latent vector $z_t \in \mathbb{R}^D$.

The default configuration uses latent dimension $D = 192$, progress dimension $k = 32$, three graph message-passing layers, one temporal recurrent layer, and dropout 0.1. These values are chosen as a conservative first configuration for a PyTorch environment. The executed smoke run in this repository uses a smaller CPU configuration to keep the experiment reproducible on the current machine.

3.3 Progress and Content Subspaces

LeWM-SDN splits the latent state into progress and content parts:

$$z_t = [z_{\text{prog},t}, z_{\text{cont},t}], \quad z_{\text{prog},t} \in \mathbb{R}^k, \quad z_{\text{cont},t} \in \mathbb{R}^{D-k}. \quad (6)$$

The progress subspace is intended to preserve temporal phase or rhythm. The content subspace stores instantaneous network state. The first two progress dimensions define a phase compass:

$$\theta_t = \text{atan2}(z_{\text{prog},t}^{(2)}, z_{\text{prog},t}^{(1)}). \quad (7)$$

The wrapped phase drift is:

$$\Delta\theta_t = |\text{atan2}(\sin(\theta_t - \theta_{t-1}), \cos(\theta_t - \theta_{t-1}))|. \quad (8)$$

3.4 Prediction Objective and SIGReg

The predictor receives latent states and action embeddings and predicts the next latent state:

$$\hat{z}_{t+1} = f_\phi(z_{t-H+1:t}, u_{t-H+1:t}). \quad (9)$$

The primary loss is next-latent mean squared error:

$$\mathcal{L}_{pred} = \frac{1}{BTD} \sum_{b,t} \|\hat{z}_{b,t+1} - z_{b,t+1}\|_2^2. \quad (10)$$

SIGReg is applied only to z_{cont} . Given random unit projections a_p and characteristic-function probe points τ_q , the implementation approximates an Epps–Pulley-style Gaussianity statistic:

$$\text{SIGReg}(Z) = \frac{1}{P} \sum_{p=1}^P \sum_q w_q \left[\left(\mathbb{E} \cos(\tau_q a_p^\top Z) - e^{-\tau_q^2/2} \right)^2 + \left(\mathbb{E} \sin(\tau_q a_p^\top Z) \right)^2 \right]. \quad (11)$$

Restricting SIGReg to z_{cont} avoids imposing a Gaussian prior on the phase-like progress subspace.

The full training objective is:

$$\mathcal{L} = \mathcal{L}_{pred} + \lambda_s \text{SIGReg}(z_{\text{cont}}) + \lambda_r \mathcal{L}_{trip}(z_{\text{prog}}) + \lambda_{st} \mathcal{L}_{straight}(z). \quad (12)$$

The triplet term encourages adjacent progress states to be closer than states further apart, and the straightening term encourages consecutive latent velocities to align:

$$\mathcal{L}_{straight} = -\mathbb{E}_t \cos(z_{t+1} - z_t, z_{t+2} - z_{t+1}). \quad (13)$$

3.5 Anomaly Score

The deployed detector scores each transition with prediction surprise and phase drift:

$$A_t = \alpha \|\hat{z}_{t+1} - z_{t+1}\|_2^2 + \beta \Delta\theta_t. \quad (14)$$

Thresholds are selected from normal validation windows, e.g. the 99th percentile of normal scores. This is an anomaly-detection setup: the predictive model can be trained only on normal windows, unlike supervised classifiers that require attack labels.

3.6 Mitigation Planning Hook

The repository also contains a latent rollout and defense-cost hook for future MPC/CEM mitigation. Given candidate action sequences $u_{t:t+H-1}^{(s)}$, the model rolls out latent futures $\hat{z}_{t+1:t+H}^{(s)}$ and evaluates:

$$C^{(s)} = \sum_{j=1}^H \left[\eta_c \text{Congestion}(\hat{z}_{t+j}^{(s)}) + \eta_l \text{PacketLoss}(\hat{z}_{t+j}^{(s)}) + \eta_u \text{Cost}(u_{t+j-1}^{(s)}) \right]. \quad (15)$$

Only the first action of the best sequence would be executed before observing the next network state and replanning. This paper does not claim mitigation results yet because public intrusion datasets do not contain logged mitigation actions.

4 Implementation and Reproducibility

The repository is organized around an SDN-first pipeline. The main code paths are:

- `module.py`: graph message passing, temporal graph encoder, mitigation action encoder, transformer predictor, and SIGReg.
- `jepa.py`: SDN-JEPA model, latent phase, anomaly score, rollout, and defense cost.
- `sdn_data.py`: graph-window dataset loader for `.npz`, `.pt`, and `.pth` files.
- `train.py`: PyTorch training loop for the temporal GNN JEPA.
- `eval.py`: anomaly-score evaluation for trained checkpoints.
- `scripts/generate_synthetic_sdn.py`: offline SDN/DDoS graph generator.
- `scripts/run_baseline_comparison.py`: offline baseline comparison.
- `scripts/prepare_insdn.py` and `scripts/prepare_unsw_nb15.py`: public-dataset conversion scripts.

The current machine has NumPy, PIL, and a local PyTorch virtual environment, but no LaTeX compiler. Therefore, the PyTorch training path is executed as a small CPU smoke experiment, while NumPy experiments provide faster baseline comparison, figures, and report generation.

5 Datasets

5.1 Preferred Public Benchmarks

The preferred public benchmarks are InSDN and UNSW-NB15. InSDN is the closest match to the SDN problem because it was constructed for SDN intrusion detection. UNSW-NB15 is smaller and easier to copy offline as CSV, making it a practical fallback. CICIDS2017 and CICDDoS2019 are useful DDoS-oriented benchmarks but are larger and less convenient under restricted access.

Table 1. Converted public NF-UNSW-NB15 graph dataset.

Property	Value
Original flows	1,623,118
Graph timesteps	1,624
Nodes	44
Directed edges	294
Node feature tensor	(1624, 44, 8)
Edge feature tensor	(1624, 294, 5)
Normal timesteps	540
Attack timesteps	1,084

The repository now contains conversion scripts for InSDN and UNSW-NB15. The InSDN converter maps flow records to endpoint graphs when IP fields are available. The UNSW-NB15 converter handles the common public train/test CSVs, which may not include raw IP addresses, by building a fixed heterogeneous graph over protocol, service, state, and global traffic summary nodes. This is less topologically faithful than InSDN, but it provides a reproducible public CSV path once files are copied into the workspace.

5.2 Public NF-UNSW-NB15 Benchmark

To remove the purely synthetic-data limitation, the revised repository also downloads and converts the public NF-UNSW-NB15 NetFlow dataset [7]. This dataset is a public NetFlow representation of UNSW-NB15 with endpoint IP addresses, ports, protocols, packet/byte counters, flow duration, a binary label, and attack name. The downloaded archive contains `NF-UNSW-NB15.csv`. The converter groups every 1,000 flows into one graph snapshot and keeps the top active endpoints as graph nodes.

5.3 Offline Controlled Dataset

The executed experiment still includes a generated SDN/DDoS graph dataset as a controlled functional test. The generator creates one controller node, switch nodes, host nodes, directed topology edges, normal periodic traffic, and bursty DDoS periods. The dataset is not a final benchmark; it is used to validate the mechanics of the proposed pipeline under a known topology.

6 Experiments

6.1 Compared Methods

The expert feedback correctly noted that reporting only one model is not sufficient. The revised experiment therefore compares ten methods:

Synthetic SDN/DDoS Timeline

1200 timesteps: normal windows followed by periodic DDoS bursts



Fig. 2. Timeline of the offline SDN/DDoS dataset. Normal traffic occupies the initial and inter-attack periods; DDoS bursts appear periodically after the normal prefix.

Table 2. Offline synthetic SDN/DDoS graph dataset.

Property	Value
Timesteps	1,200
Nodes	29
Directed edges	64
Node feature tensor (1200, 29, 8)	
Edge feature tensor (1200, 64, 5)	
Normal timesteps	960
Attack timesteps	240

1. **One-class feature distance:** computes the squared distance from the mean normal graph feature vector and thresholds the 99th percentile of normal scores.
2. **Supervised logistic regression:** a NumPy binary classifier trained on labeled warmup windows.
3. **Supervised shallow MLP:** a one-hidden-layer NumPy neural classifier trained on the same labeled warmup windows.
4. **Supervised gradient-boostered stumps:** a lightweight XGBoost-style boosted decision-stump classifier implemented in NumPy because the current environment does not include the external XGBoost package.
5. **Supervised PyTorch MLP:** a supervised neural classifier trained on fixed-length graph-feature windows.
6. **Supervised PyTorch CNN:** a one-dimensional convolutional classifier over graph-feature windows.
7. **Supervised PyTorch LSTM:** a recurrent sequence classifier over graph-feature windows.
8. **Raw-feature ridge surprise:** a next-state predictor in raw graph-feature space trained on normal windows.
9. **LeWM-SDN latent surprise + phase:** the proposed reference detector using normal-state PCA latent encoding, next-latent ridge prediction, and phase drift.

10. **PyTorch temporal GNN JEPA**: the actual neural SDN-JEPA architecture trained for a short CPU smoke run on normal synthetic windows.

The supervised baselines use attack labels during training and are therefore not directly equivalent to the anomaly-detection setup. They are included to answer the reviewer question “compared with what?” and to expose whether the offline dataset is too easy.

6.2 Protocol

The exact commands executed were:

```
python scripts/generate_synthetic_sdn.py \  
  --output data/processed/synthetic_sdn_ddos.npz \  
  --steps 1200 \  
  --seed 3072  
  
python scripts/run_baseline_comparison.py \  
  --data data/processed/synthetic_sdn_ddos.npz \  
  --output outputs/eval/baseline_comparison.json \  
  --history 8 \  
  --latent-dim 16 \  
  --warmup 600  
  
python scripts/run_torch_supervised_baselines.py \  
  --data data/processed/synthetic_sdn_ddos.npz \  
  --output outputs/eval/torch_supervised_baselines.json \  
  --history 8 \  
  --warmup 600 \  
  --hidden-dim 64 \  
  --epochs 80  
  
python train.py \  
  data.path=data/processed/synthetic_sdn_ddos.npz \  
  data.num_steps=10 history_size=4 embed_dim=32 prog_dim=8 \  
  model.encoder.hidden_dim=32 model.encoder.depth=2 \  
  model.predictor.depth=1 model.predictor.heads=4 \  
  model.predictor.mlp_dim=128 model.predictor.dim_head=8 \  
  trainer.max_epochs=8 loader.batch_size=64 \  
  output_dir=outputs/checkpoints/synthetic_torch \  
  output_model_name=synthetic_torch loss.sigreg.kwargs.num_proj=16  
  
python eval.py \  
  checkpoint=outputs/checkpoints/synthetic_torch/synthetic_torch_best.pt \  
  data.path=data/processed/synthetic_sdn_ddos.npz \  
  data.num_steps=10 output.path=outputs/eval/torch_sdn_jepa_eval.json
```

Baseline Comparison on Offline SDN/DDoS Dataset




Fig. 3. F1 comparison between the proposed reference detector and baselines on the offline SDN/DDoS dataset. Supervised logistic regression and neural baselines achieve very high F1, indicating that the synthetic dataset is easy for label-supervised classifiers.

The first 600 timesteps are used as the warmup/training region for supervised baselines. Anomaly methods estimate their thresholds from normal scores. NumPy methods and supervised PyTorch baselines are evaluated on aligned post-history transitions. The temporal GNN JEPA run uses overlapping graph windows from the same synthetic dataset and is reported as a short CPU training smoke test rather than a tuned benchmark run.

For NF-UNSW-NB15, the source CSV is sorted with attack-heavy regions before normal-heavy regions. A prefix split would therefore be misleading. The public benchmark uses a stratified 50/50 train/test split over graph windows. The normal-only anomaly thresholds are estimated from the training normals, and reported metrics are computed on the held-out test split.

7 Results

Table 4 reports the public-data benchmark run. It shows a different picture from the synthetic result. Supervised classifiers perform well on NF-UNSW-NB15, while the current one-class and predictive surprise variants have poor recall. The Temporal GNN JEPA public smoke run trains successfully but obtains $F1 = 0$ under the default normal-score threshold; an oracle threshold diagnostic reaches $F1 = 0.8010$, which indicates that the score distribution contains label information but that threshold calibration is currently inadequate. This public result is negative for the current anomaly formulation but valuable: it prevents the paper from relying only on an easy generated dataset.

LeWM-SDN Reference Anomaly Score Separation

Attack windows have much higher predictive latent surprise.




Fig. 4. Mean anomaly score separation for the LeWM-SDN reference detector. DDoS windows have substantially higher predictive latent surprise than normal windows.

The LeWM-SDN reference detector reaches high recall and F1, but it does not outperform the supervised logistic, boosting, CNN, LSTM, or MLP classifiers on this synthetic dataset. This is an important correction to the earlier draft: the current evidence demonstrates that the predictive latent pipeline runs and is competitive with unsupervised feature-distance detection, not that it is superior to supervised ML/DL on a public benchmark. The perfect supervised scores also suggest that the generated dataset is too separable to support strong scientific claims.

The PyTorch temporal GNN JEPA now has an executed training/evaluation trace: training loss decreases from 4.8049 to 3.7158 over eight CPU epochs and validation loss decreases from 3.7971 to 3.1776. However, its anomaly F1 is only 0.0062 under the default 99th-percentile normal threshold. This is a negative but useful result: the neural architecture is runnable, yet the current small synthetic setup and short CPU training are not sufficient to claim neural detection quality.

Nevertheless, the result is useful for engineering validation. The raw-feature next-state baseline performs poorly, while the latent surprise + phase version works well. This supports the claim that latent representation and phase-aware scoring are useful even in the controlled setting.

8 Discussion

The revised experiment directly addresses the absence of baselines. It also shows why the synthetic dataset cannot be the final benchmark: supervised classifiers solve it perfectly. A publishable Q4-level paper should therefore run the PyTorch temporal GNN JEPA on the public NF-UNSW-NB15 graph or on InSDN and compare against classical ML and DL baselines. The current public benchmark

Table 3. Baseline comparison on the offline SDN/DDoS dataset.

Method	Precision	Recall	F1	Accuracy
One-class feature distance	0.9449	1.0000	0.9717	0.9883
Supervised logistic regression	1.0000	1.0000	1.0000	1.0000
Supervised NumPy MLP	1.0000	1.0000	1.0000	1.0000
Supervised boosted stumps	0.9796	1.0000	0.9897	0.9958
Supervised PyTorch MLP	0.9875	0.9875	0.9875	0.9950
Supervised PyTorch CNN	0.9833	0.9833	0.9833	0.9933
Supervised PyTorch LSTM	0.9833	0.9833	0.9833	0.9933
Raw-feature ridge surprise	0.7091	0.1625	0.2644	0.8180
LeWM-SDN latent surprise + phase	0.9057	1.0000	0.9505	0.9790
PyTorch temporal GNN JEPA (8 epochs)	0.0753	0.0032	0.0062	0.7911

Table 4. Baseline comparison on the public NF-UNSW-NB15 graph dataset using a stratified held-out test split.

Method	Precision	Recall	F1	Accuracy
One-class feature distance	0.8421	0.0595	0.1111	0.3663
Supervised logistic regression	0.9818	1.0000	0.9908	0.9876
Supervised NumPy MLP	0.9818	1.0000	0.9908	0.9876
Supervised boosted stumps	0.9809	0.9554	0.9680	0.9579
Supervised PyTorch MLP	0.9779	0.9888	0.9834	0.9777
Supervised PyTorch CNN	0.9781	0.9963	0.9871	0.9827
Supervised PyTorch LSTM	0.9799	0.9963	0.9880	0.9839
Raw-feature ridge surprise	0.8000	0.0669	0.1235	0.3676
LeWM-SDN latent surprise + phase	0.7273	0.0446	0.0841	0.3527
PyTorch temporal GNN JEPA (8 epochs)	0.0000	0.0000	0.0000	0.3285

already includes logistic regression, a lightweight gradient-boosted stump baseline, MLP, one-class distance, and next-state surprise baselines. The synthetic benchmark also includes PyTorch CNN and LSTM baselines. The next public benchmark iteration should add external packages such as XGBoost or random forest where dependencies are available.

The feedback also noted that the original contribution claims were broader than the experiments. This revision narrows the claims. We do not claim that MPC/CEM mitigation has been validated, only that the model exposes a latent rollout and defense-cost hook. We also do not claim that the temporal GNN JEPA is competitive yet. The implemented PyTorch training path now runs, but it requires public data and tuning before it can support a publication claim.

Public NF-UNSW-NB15 Test F1 Comparison




Fig. 5. Held-out test F1 comparison on the public NF-UNSW-NB15 graph dataset. Supervised baselines perform well, while the current normal-only surprise scores and the public Temporal GNN JEPA smoke run are not competitive.

9 Threats to Validity

Synthetic data. The generated SDN/DDoS dataset is much smaller than public intrusion datasets. Its result should not be treated as benchmark evidence. It validates executability, code integration, and the anomaly-score mechanism. The revised paper therefore also includes a public NF-UNSW-NB15 graph experiment.

Reference and neural implementations. The strongest executed LeWM-SDN result still comes from the PCA-ridge NumPy reference. The PyTorch temporal GNN JEPA is now trained and evaluated, but its detection score is weak in the current short CPU run. This validates executability, not final neural performance.

Baseline fairness. Supervised classifiers receive labels during training, whereas the anomaly detectors estimate normal behavior. This makes the supervised methods stronger on both the synthetic and public split. The public NF-UNSW-NB15 split is stratified because the source CSV is not temporally randomized.

Missing real mitigation actions. Public intrusion datasets usually do not log SDN mitigation actions. The MPC/CEM part requires a controlled SDN testbed such as Mininet with Ryu or ONOS, where actions and post-action states can be recorded.

10 Conclusion

LeWM-SDN converts the world-model idea from pixel-based physical environments to graph-based SDN network dynamics. The revised project now uses Springer LNCS formatting, includes related work, gives explicit method equations, adds figures, compares against baselines, and provides public-dataset conversion scripts. The current executable result remains an offline controlled validation rather than a final public benchmark. The next decisive step is to run the implemented temporal GNN JEPA on the public NF-UNSW-NB15 graph or InSDN, tune it beyond the current CPU smoke run, and compare it with the existing boosting/CNN/LSTM/MLP baselines plus external random-forest, XGBoost, and autoencoder baselines.

References

1. McKeown, N., Anderson, T., Balakrishnan, H., Parulkar, G., Peterson, L., Rexford, J., Shenker, S., Turner, J.: OpenFlow: enabling innovation in campus networks. *ACM SIGCOMM Computer Communication Review* 38(2), 69–74 (2008)
2. Kreutz, D., Ramos, F.M.V., Verissimo, P.E., Rothenberg, C.E., Azodolmolky, S., Uhlig, S.: Software-defined networking: A comprehensive survey. *Proceedings of the IEEE* 103(1), 14–76 (2015)
3. Yan, Q., Yu, F.R., Gong, Q., Li, J.: Software-defined networking and distributed denial of service attacks in cloud computing environments: A survey. *IEEE Communications Surveys & Tutorials* 18(1), 602–622 (2016)
4. Scott-Hayward, S., O’Callaghan, G., Sezer, S.: SDN security: A survey. *IEEE SDN for Future Networks and Services* (2013)
5. Tavallaee, M., Bagheri, E., Lu, W., Ghorbani, A.A.: A detailed analysis of the KDD CUP 99 data set. *IEEE Symposium on Computational Intelligence for Security and Defense Applications* (2009)
6. Moustafa, N., Slay, J.: UNSW-NB15: a comprehensive data set for network intrusion detection systems. *Military Communications and Information Systems Conference* (2015)
7. Sarhan, M., Layeghy, S., Moustafa, N., Portmann, M.: NetFlow datasets for machine learning-based network intrusion detection systems. *Big Data Technologies and Applications* (2021)
8. Sharafaldin, I., Lashkari, A.H., Ghorbani, A.A.: Toward generating a new intrusion detection dataset and intrusion traffic characterization. *ICISSP* (2018)
9. Sharafaldin, I., Lashkari, A.H., Hakak, S., Ghorbani, A.A.: Developing realistic distributed denial of service dataset and taxonomy. *ICCST* (2019)
10. Koroniotis, N., Moustafa, N., Sitnikova, E., Turnbull, B.: Towards the development of realistic botnet dataset in the Internet of Things for network forensic analytics: Bot-IoT dataset. *Future Generation Computer Systems* 100, 779–796 (2019)
11. Moustafa, N.: A new distributed architecture for evaluating AI-based security systems at the edge: Network TON_IoT datasets. *Sustainable Cities and Society* 72 (2021)
12. Elsayed, M.S., Le-Khac, N.A., Dev, S., Jurcut, A.D.: InSDN: A novel SDN intrusion dataset. *IEEE Access* 8, 165263–165284 (2020)

13. Kim, J., Kim, J., Thu, H.L.T., Kim, H.: Long short term memory recurrent neural network classifier for intrusion detection. *International Conference on Platform Technology and Service* (2016)
14. Vinayakumar, R., Alazab, M., Soman, K.P., Poornachandran, P., Al-Nemrat, A., Venkatraman, S.: Deep learning approach for intelligent intrusion detection system. *IEEE Access* 7, 41525–41550 (2019)
15. Javaid, A., Niyaz, Q., Sun, W., Alam, M.: A deep learning approach for network intrusion detection system. *EAI International Conference on Bio-inspired Information and Communications Technologies* (2016)
16. Mirsky, Y., Doitshman, T., Elovici, Y., Shabtai, A.: Kitsune: an ensemble of autoencoders for online network intrusion detection. *NDSS* (2018)
17. Kipf, T.N., Welling, M.: Semi-supervised classification with graph convolutional networks. *International Conference on Learning Representations* (2017)
18. Velickovic, P., Cucurull, G., Casanova, A., Romero, A., Lio, P., Bengio, Y.: Graph attention networks. *International Conference on Learning Representations* (2018)
19. Wu, Z., Pan, S., Chen, F., Long, G., Zhang, C., Yu, P.S.: A comprehensive survey on graph neural networks. *IEEE Transactions on Neural Networks and Learning Systems* 32(1), 4–24 (2021)
20. Zhang, X., Liu, J., Yang, H., Chen, Y.: Graph neural networks for network security: A survey. *Computers & Security* 122 (2022)
21. Ha, D., Schmidhuber, J.: World models. *arXiv preprint arXiv:1803.10122* (2018)
22. Hafner, D., Lillicrap, T., Fischer, I., Villegas, R., Ha, D., Lee, H., Davidson, J.: Learning latent dynamics for planning from pixels. *ICML* (2019)
23. Hafner, D., Lillicrap, T., Ba, J., Norouzi, M.: Dream to control: Learning behaviors by latent imagination. *International Conference on Learning Representations* (2020)
24. LeCun, Y.: A path towards autonomous machine intelligence. *Open manuscript* (2022)
25. Bardes, A., Garrido, Q., Ponce, J., LeCun, Y., Rabbat, M., Assran, M., Ballas, N.: Revisiting feature prediction for learning visual representations from video. *arXiv preprint* (2024)
26. Maes, L., Le Lidec, Q., Scieur, D., LeCun, Y., Balestrieri, R.: LeWorldModel: Stable end-to-end joint-embedding predictive architecture from pixels. *arXiv preprint* (2026)
27. Chandola, V., Banerjee, A., Kumar, V.: Anomaly detection: A survey. *ACM Computing Surveys* 41(3), 1–58 (2009)
28. Ahmad, S., Lavin, A., Purdy, S., Agha, Z.: Unsupervised real-time anomaly detection for streaming data. *Neurocomputing* 262, 134–147 (2017)
29. Malhotra, P., Vig, L., Shroff, G., Agarwal, P.: Long short term memory networks for anomaly detection in time series. *ESANN* (2015)